

```
In [ ]: import pandas as pd
import numpy as np
import altair as alt

import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
warnings.simplefilter(action='ignore', category=UserWarning)

from sklearn.preprocessing import QuantileTransformer
from sklearn.metrics import accuracy_score, classification_report
from sklearn.model_selection import train_test_split
```

```
In [ ]: df = pd.read_csv("./assets/diabetes.csv")
```

```
In [ ]: df.head(5)
```

```
Out[ ]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

```
In [ ]: df.describe()
```

```
Out[ ]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	0.471876	33.240885	0.348958
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.078000	21.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	0.243750	24.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	0.372500	29.000000	0.000000
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	0.626250	41.000000	1.000000
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	2.420000	81.000000	1.000000

```
In [ ]: df.corr("pearson")
```

```
Out[ ]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
Pregnancies	1.000000	0.129459	0.141282	-0.081672	-0.073535	0.017683	-0.033523	0.544341	0.221898
Glucose	0.129459	1.000000	0.152590	0.057328	0.331357	0.221071	0.137337	0.263514	0.466581
BloodPressure	0.141282	0.152590	1.000000	0.207371	0.088933	0.281805	0.041265	0.239528	0.065068
SkinThickness	-0.081672	0.057328	0.207371	1.000000	0.436783	0.392573	0.183928	-0.113970	0.074752
Insulin	-0.073535	0.331357	0.088933	0.436783	1.000000	0.197859	0.185071	-0.042163	0.130548
BMI	0.017683	0.221071	0.281805	0.392573	0.197859	1.000000	0.140647	0.036242	0.292695
DiabetesPedigreeFunction	-0.033523	0.137337	0.041265	0.183928	0.185071	0.140647	1.000000	0.033561	0.173844
Age	0.544341	0.263514	0.239528	-0.113970	-0.042163	0.036242	0.033561	1.000000	0.238356
Outcome	0.221898	0.466581	0.065068	0.074752	0.130548	0.292695	0.173844	0.238356	1.000000

```
In [ ]: cor_data = df.corr("pearson").stack().reset_index().rename(columns={0: 'correlation', 'level_0': 'variable', 'level_1': 'variable2'})
```


```
base = alt.Chart(cor_data).encode(
    x='variable2:0',
    y='variable:0'
)

text = base.mark_text().encode(
    text='correlation_label',
    color=alt.condition(
        alt.datum.correlation > 0.5,
        alt.value('white'),
        alt.value('black')
    )
)

cor_plot = base.mark_rect().encode(
    color='correlation:Q'
)

cor_plot
```

Out[]:



```
In [ ]: df['Glucose'] = df['Glucose'].replace(0, df['Glucose'].mean())
# Correcting missing values in blood pressure
df['BloodPressure'] = df['BloodPressure'].replace(0, df['BloodPressure'].mean()) # There are 35 records with 0 BloodPressure in data
# Correcting missing values in BMI
df['BMI'] = df['BMI'].replace(0, df['BMI'].median())
# Correct missing values in Insulin and SkinThickness

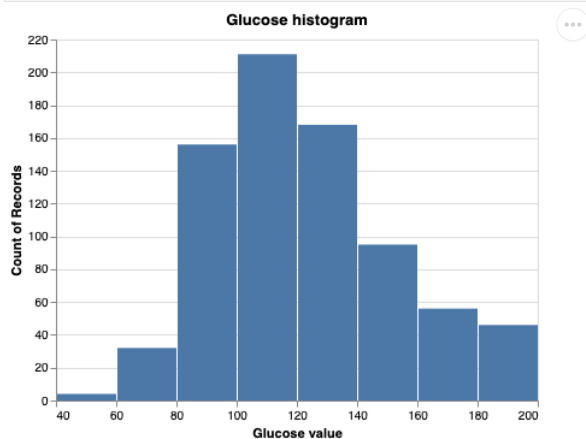
df['SkinThickness'] = df['SkinThickness'].replace(0, df['SkinThickness'].median())
df['Insulin'] = df['Insulin'].replace(0, df['Insulin'].median())

df.describe()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	121.681605	72.254807	27.334635	94.652344	32.450911	0.471876	33.240885	0.348958
std	3.369578	30.436016	12.115932	9.229014	105.547598	6.875366	0.331329	11.760232	0.476951
min	0.000000	44.000000	24.000000	7.000000	14.000000	18.200000	0.078000	21.000000	0.000000
25%	1.000000	99.750000	64.000000	23.000000	30.500000	27.500000	0.243750	24.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	31.250000	32.000000	0.372500	29.000000	0.000000
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	0.626250	41.000000	1.000000
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	2.420000	81.000000	1.000000

```
In [ ]: alt.Chart(df, title="Glucose histogram").mark_bar().encode(
    x= alt.X("Glucose:Q", bin=True, title="Glucose value"),
    y='count()',
)
```

Out[]:



```
In [ ]: # Regression SkinThickness/BMI
chart = alt.Chart(df).mark_point().encode(
    x= alt.X("SkinThickness"),
    y= alt.Y("BMI"),
)

(chart + chart.transform_regression('SkinThickness', 'BMI').mark_line().encode(
    color=alt.value("#FFAA00")
)).display()

# Regression BloodPressure/Age
chart2 = alt.Chart(df).mark_point().encode(
    x= alt.X("BloodPressure"),
    y= alt.Y("Age"),
)

(chart2 + chart2.transform_regression('BloodPressure', 'Age').mark_line().encode(
    color=alt.value("#FFAA00")
)).display()




# Regression Glucose/DiabetesPedigreeFunction
chart3 = alt.Chart(df).mark_point().encode(
    x= alt.X("Glucose"),
```

```

    y= alt.Y("DiabetesPedigreeFunction")
)

(chart3 + chart3.transform_regression('Glucose', 'DiabetesPedigreeFunction').mark_line().encode(
    color=alt.value("#FFAA00")
)).display()

```



```

In [ ]: diabetes_women = pd.DataFrame({
    'pregnancies': df["Pregnancies"][df["Outcome"] == 1].value_counts().to_frame().index.to_list(),
    'frequency': df["Pregnancies"][df["Outcome"] == 1].value_counts().values,
    'name': "diabetes"
})

non_diabetes_women = pd.DataFrame({
    'pregnancies': df["Pregnancies"][df["Outcome"] == 0].value_counts().to_frame().index.to_list(),
    'frequency': df["Pregnancies"][df["Outcome"] == 0].value_counts().values,
    'name': "non diabetes"
})

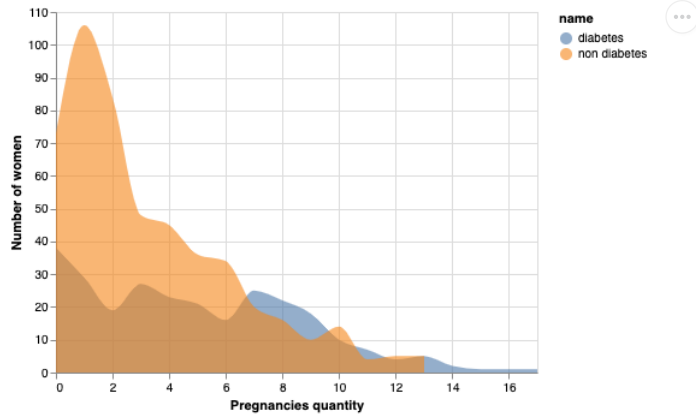
area1 = alt.Chart(diabetes_women).mark_area(
    interpolate='monotone'
).encode(
    alt.X("pregnancies", title="Pregnancies quantity"),
    alt.Y("frequency", title='Number of women'),
    opacity=alt.value(0.6),
    color="name"
)

area2 = alt.Chart(non_diabetes_women).mark_area(
    interpolate='monotone'
).encode(
    alt.X("pregnancies", scale=alt.Scale(zero=False, nice=False), title="Pregnancies quantity"),
    alt.Y("frequency", title='Number of women'),
    opacity=alt.value(0.6),
    color="name"
)

```

area1 + area2

Out []:




In []:

```
source = pd.DataFrame({
    'a': ['diabetes', 'non-diabetes'],
    'b': [df[df["Outcome"] == 1]["Outcome"].count(), df[df["Outcome"] == 0]["Outcome"].count()]
})

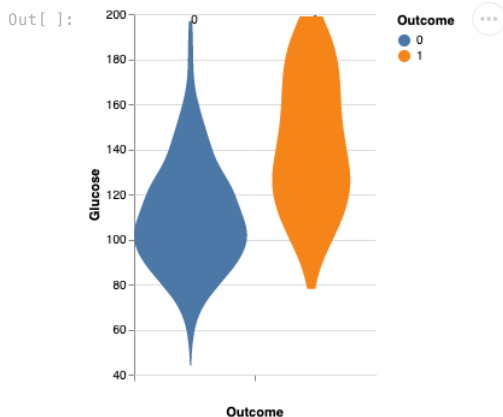
alt.Chart(source).mark_bar().encode(
    alt.X("a", title=""),
    alt.Y("b", title='Number of people in sample'),
).properties(
    width=200,
    height=300
)
```

Out []:



In []:

```
alt.Chart(df).transform_density(
    'Glucose',
    as_=['Glucose', 'density'],
    groupby=['Outcome']
).mark_area(orient='horizontal').encode(
    y='Glucose:Q',
    color='Outcome:N',
    x=alt.X(
        'density:Q',
        stack='center',
        impute=None,
        title=None,
        axis=alt.Axis(labels=False, values=[0], grid=False, ticks=True),
    ),
    column=alt.Column(
        'Outcome:N',
        header=alt.Header(
            titleOrient='bottom',
            labelOrient='bottom',
            labelPadding=0,
        ),
    ),
).properties(
    width=100
).configure_facet(
    spacing=0
).configure_view(
    stroke=None
)
```



```
In [ ]: # Data Transformation
q = QuantileTransformer()
X = q.fit_transform(df)
transformedDF = q.transform(X)
transformedDF = pd.DataFrame(X)
transformedDF.columns = ['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome']
transformedDF.head()
```

Out[]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	0.747718	0.810300	0.516949	0.801825	0.256193	0.591265	0.750978	0.889831	1.0
1	0.232725	0.091265	0.290091	0.644720	0.256193	0.213168	0.475880	0.558670	0.0
2	0.863755	0.956975	0.233377	0.357888	0.256193	0.077575	0.782269	0.585398	1.0
3	0.232725	0.124511	0.290091	0.357888	0.662973	0.284224	0.106258	0.000000	0.0
4	0.000000	0.721643	0.005215	0.801825	0.834420	0.926988	0.997392	0.606258	1.0

```
In [ ]: features = transformedDF.drop(["Outcome"], axis=1)
labels = transformedDF["Outcome"]
# train (70%) and test (30%)
x_train, x_test, y_train, y_test = train_test_split(features, labels, test_size=0.30, random_state=7)
print (f'Shape of Train Data : {x_train.shape}')
print (f'Shape of Test Data : {x_test.shape}')
```

Shape of Train Data : (537, 8)
Shape of Test Data : (231, 8)

Validation data will be also useful to provide an unbiased **evaluation of a model** fit on the **training dataset** while tuning model hyperparameters. Model occasionally sees this data, but never does it *Learn* from this. Validation set is also called dev set.



```
In [ ]: x_dev, x_dev_test, y_dev, y_dev_test = train_test_split(x_test, y_test, test_size=0.5, random_state=2)
print (f'Shape of Dev/Validation Data : {x_dev.shape}')
```

Shape of Dev/Validation Data : (115, 8)

We've prepared training, test and validation data. Question is how many hidden layers and neurons should we use? This is common question and also if we googled internet in appropriate way ... this is impossible to have only one answer. In **Introduction to Neural Networks for Java**, Second Edition by *jeffheaton* we've found definition:

Table 5.1: Determining the Number of Hidden Layers

Number of Hidden Layers	Result
none	Only capable of representing linear separable functions or decisions.
1	Can approximate any function that contains a continuous mapping from one finite space to another.
2	Can represent an arbitrary decision boundary to arbitrary accuracy with rational activation functions and can approximate any smooth mapping to any accuracy.

Base on our case two hidden layers will be good choice. There is still an issue how many neurons per layer should we use.

1. For the first input layer should be 8 neurons because we have 8 features/columns.
2. Output layer should contain 1 neuron because our result is 1 or 0 (diabetic or not).

Using too few neurons in the hidden layers will result in something called underfitting. Underfitting occurs when there are too few neurons in the hidden layers to adequately detect the signals in a complicated data set.

Using too many neurons in the hidden layers can result in several problems. First, too many neurons in the hidden layers may result in overfitting. Overfitting occurs when the neural network has so much information processing capacity that the limited amount of information contained in the training set is not enough to train all of the neurons in the hidden layers.

We decided to use this formula:

$$N_h = \frac{N_s}{a * (N_i + N_o)}$$

N_i = number of input neurons.

N_o = number of output neurons.

N_s = number of samples in training data set.

a = an arbitrary scaling factor usually 2-10.


Also with formula above we wanted to follow rules:

- The number of hidden neurons should be between the size of the input layer and the size of the output layer. $\langle 1; 8 \rangle$
- The number of hidden neurons should be $2/3$ the size of the input layer, plus the size of the output layer. $8 * 2/3 + 1 = 6.3$
- The number of hidden neurons should be less than twice the size of the input layer. $x < 2 * 8 \Rightarrow x < 16$

So we decided:

$$N_{h1} = \frac{537}{5 * (8 + 1)} = 12$$

$$N_{h1} = \frac{537}{10 * (8 + 1)} = 6$$



```
In [ ]: from tensorflow.keras.models import Sequential, Model
from tensorflow.keras.layers import Activation, Dense, Dropout, BatchNormalization, Input
from tensorflow.keras.optimizers import Adam

def create_model_and_fit(epochs):
    print(f"START CREATING MODEL FOR EPOCH={epochs}")
    inputs = Input(name='inputs', shape=[x_train.shape[1],])
    layer = Dense(12, name='FC1')(inputs)
    layer = BatchNormalization(name='BC1')(layer)
    layer = Activation('relu', name='Activation1')(layer)
    layer = Dropout(0.3, name='Dropout1')(layer)
    layer = Dense(6, name='FC2')(layer)
    layer = BatchNormalization(name='BC2')(layer)
    layer = Activation('relu', name='Activation2')(layer)
    layer = Dropout(0.3, name='Dropout2')(layer)
    layer = Dense(1, name='OutLayer')(layer)
```

```

layer = Activation('sigmoid', name='sigmoid')(layer)
model = Model(inputs=inputs, outputs=layer)

model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
historic = model.fit(x=x_train, y=y_train, epochs = epochs, verbose=0)

loss_df = pd.DataFrame({
    'loss': historic.history["loss"],
    'epoch': [i for i, x in enumerate(historic.history["loss"])],
    "name": "Error"
})

accuracy_df = pd.DataFrame({
    'accuracy': historic.history["accuracy"],
    'epoch': [i for i, x in enumerate(historic.history["accuracy"])],
    "name": "Accuracy"
})

alt1 = alt.Chart(loss_df, title=f"Binary cross entropy on train dataset, EPOCH = {epochs}").mark_line(interpolate='basis').encode(
    x = 'epoch',
    y = alt.Y('loss', title='value', scale=alt.Scale(domain=[min(historic.history["loss"]), max(historic.history["loss"])])),
    color = "name"
)
alt2 = alt.Chart(accuracy_df).mark_line(interpolate='basis').encode(
    x='epoch',
    y = alt.Y('accuracy', title='value', scale=alt.Scale(domain=[min(historic.history["accuracy"]), max(historic.history["accuracy"])])),
    color = "name"
)

(alt1 + alt2).display()

return model;

```

```

In [ ]: def show_report(x, y, name, model, epoch):
    print(f"Report for {name} set\n")
    y_prediction = model.predict(x)
    y_prediction = np.around(y_prediction)
    y_prediction = np.asarray(y_prediction)
    print('\tAccuracy:{:0.3f}\n\tClassification Report\n{}'.format(accuracy_score(y, y_prediction),
                                                                    classification_report(y, y_prediction)))

    dataResult1 = pd.DataFrame({
        'result': y,
        'name': f"y_{name}",
        'item_number': [i for i, item in enumerate(y)]
    })
    dataResult2 = pd.DataFrame({
        'result': [item[0] for item in y_prediction],
        'name': f"y_{name}_prediction",
        'item_number': [i for i, item in enumerate(y)]
    })

    chart1 = alt.Chart(dataResult1, title = f"Coverage level forecasted data with the original, EPOCH = {epoch}").mark_area(
        interpolate='monotone'
    ).encode(
        x = alt.X('item_number'),
        y = alt.Y('result', scale=alt.Scale(domain=[-0.5, 1.5])),
        color='name'
    ).properties(width=1200)
    chart2 = alt.Chart(dataResult2).mark_area(
        interpolate='monotone'
    ).encode(
        x = alt.X('item_number'),
        y = alt.Y('result', scale=alt.Scale(domain=[-0.5, 1.5])),
        opacity=alt.value(0.6),
        color='name'
    ).properties(width=1200)

    (chart1 + chart2).display()

```

```


In [ ]: # Test for different epoch [50, 200, 500]

epoch_case = [50, 200, 500]

for epoch in epoch_case:
    model = create_model_and_fit(epoch)
    show_report(x_train, y_train, "train", model, epoch)
    show_report(x_test, y_test, "test", model, epoch)

```


START CREATING MODEL FOR EPOCH=50



Report for train set

17/17 [=====] - 0s 2ms/step
 Accuracy:0.767
 Classification Report


	precision	recall	f1-score	support
0.0	0.83	0.81	0.82	353
1.0	0.65	0.68	0.67	184
accuracy			0.77	537
macro avg	0.74	0.75	0.74	537
weighted avg	0.77	0.77	0.77	537




Report for test set

8/8 [=====] - 0s 1ms/step
 Accuracy:0.758
 Classification Report

	precision	recall	f1-score	support
0.0	0.83	0.78	0.80	147
1.0	0.65	0.73	0.69	84
accuracy			0.76	231
macro avg	0.74	0.75	0.74	231
weighted avg	0.77	0.76	0.76	231



START CREATING MODEL FOR EPOCH=200

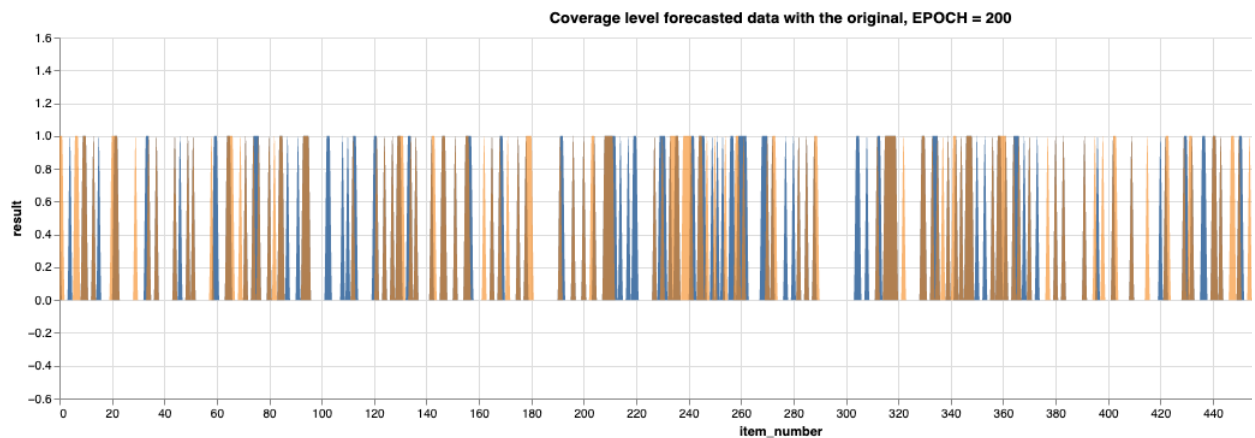


Report for train set

```
17/17 [=====] - 0s 1ms/step
Accuracy:0.788
Classification Report
precision    recall  f1-score   support

   0.0      0.82    0.87    0.84     353
   1.0      0.71    0.64    0.67     184

 accuracy          0.79     537
 macro avg         0.77     537
weighted avg         0.78     537
```




Report for test set


```
8/8 [=====] - 0s 1ms/step
Accuracy:0.762
Classification Report
precision    recall  f1-score   support

   0.0      0.80    0.84    0.82     147
   1.0      0.69    0.63    0.66      84

 accuracy          0.76     231
 macro avg         0.74     231
weighted avg         0.76     231
```



START CREATING MODEL FOR EPOCH=500




Report for train set

```
17/17 [=====] - 0s 1ms/step
Accuracy:0.814
Classification Report
precision    recall  f1-score   support

   0.0      0.86    0.86    0.86     353
   1.0      0.73    0.73    0.73     184

 accuracy          0.81     537
 macro avg         0.79    0.79    0.79     537
weighted avg         0.81    0.81    0.81     537
```



Report for test set

```
8/8 [=====] - 0s 1ms/step
Accuracy:0.766
Classification Report
precision    recall  f1-score   support

   0.0      0.81    0.82    0.82     147
   1.0      0.68    0.67    0.67      84

 accuracy          0.77     231
 macro avg         0.75    0.74    0.75     231
weighted avg         0.77    0.77    0.77     231
```

